

CS343 - Operating Systems

Module-6C

File Allocation and Free Space Management



Dr. John Jose

Assistant Professor

Department of Computer Science & Engineering

Indian Institute of Technology Guwahati, Assam.

<http://www.iitg.ac.in/johnjose/>

File-System Implementation

- ❖ File Allocation Methods
- ❖ Free-Space Management
- ❖ Efficiency and Performance
- ❖ Recovery

File-System Implementation

- ❖ **File Control Block** contains many details about the file
 - ❖ inode number, permissions, size, dates


file permissions
file dates (create, access, write)
file owner, group, ACL
file size
file data blocks or pointers to file data blocks

Contiguous Allocation

- ❖ An allocation method refers to how disk blocks are allocated for files:
- ❖ **Contiguous allocation** – each file occupies set of contiguous blocks
 - ❖ Best performance in most cases
 - ❖ Simple – only starting location (block #) and length (number of blocks) are required
 - ❖ Problems include finding space for file, knowing file size, external fragmentation, need for **compaction off-line** (**downtime**) or **on-line**

Contiguous Allocation

- ❖ Mapping from logical to physical
 - Block to be accessed = starting address
 - Displacement into block = length




Extent-Based Systems

- ❖ Many newer file systems (i.e., Veritas File System) use a modified contiguous allocation scheme
- ❖ Extent-based file systems allocate disk blocks in extents
- ❖ An **extent** is a contiguous block of disks
 - ❖ Extents are allocated for file allocation
 - ❖ A file consists of one or more extents

Linked Allocation


- ❖ **Linked allocation** – each file a linked list of blocks
 - ❖ File ends at nil pointer
 - ❖ No external fragmentation
 - ❖ Each block contains pointer to next block
 - ❖ No compaction, external fragmentation
 - ❖ Free space management system called when new block needed
 - ❖ Improve efficiency by clustering blocks into groups but increases internal fragmentation
 - ❖ Reliability can be a problem
 - ❖ Locating a block can take many I/Os and disk seeks

Linked Allocation



directory entry

test	...	217
name		start block




Indexed Allocation


❖ Indexed allocation

❖ Each file has its own **index block**(s) of pointers to its data blocks

❖ Logical view




index table



Indexed Allocation


- ❖ Need index table
- ❖ Random access
- ❖ Dynamic access without external fragmentation, but have overhead of index block
- ❖ Single Level and Multilevel Index for small and large files

Indexed Allocation



UNIX UFS

4K bytes per block, 32-bit addresses




More index blocks than can be addressed with 32-bit file pointer

Performance

- ❖ Best method depends on file access type
 - ❖ Contiguous great for sequential and random
- ❖ Linked good for sequential, not random
- ❖ Declare access type at creation → select either contiguous or linked
- ❖ Indexed more complex – multiple index block reads.

Free-Space Management

- ❖ File system maintains **free-space list** to track available blocks
- ❖ **Bit vector** or **bit map** (n blocks)



$$\text{bit}[i] = \begin{cases} 1 \Rightarrow \text{block}[i] \text{ free} \\ 0 \Rightarrow \text{block}[i] \text{ occupied} \end{cases}$$

Free-Space Management

- ❖ Bit map requires extra space

- ❖ Example:

- block size = 4KB = 2^{12} bytes

- disk size = 2^{40} bytes (1 terabyte)


- $n = 2^{40}/2^{12} = 2^{28}$ bits (or 32MB)

- ❖ Easy to get contiguous files

Linked Free Space List on Disk

❖ Linked list (free list)

- ❖ Cannot get contiguous space easily
- ❖ No waste of space
- ❖ No need to traverse the entire list (if # free blocks recorded)



Free-Space Management

❖ Grouping

- ❖ Modify linked list to store address of next $n-1$ free blocks in first free block, plus a pointer to next block that contains free-block-pointers (like this one)

❖ Counting

- ❖ Because space is frequently contiguously used and freed, with contiguous-allocation allocation, extents, or clustering
 - ❖ Keep address of first free block and count of following free blocks
 - ❖ Free space list then has entries containing addresses and counts

Free-Space Management

❖ Space Maps

- ❖ Divides device space into **metaslab** units and manages metaslabs
 - ❖ Given volume can contain hundreds of metaslabs
- ❖ Each metaslab has associated space map -uses counting algorithm

Efficiency and Performance

- ❖ Efficiency depends on:
 - ❖ Disk allocation and directory algorithms
 - ❖ Types of data kept in file's directory entry
 - ❖ Pre-allocation or as-needed allocation of metadata structures
 - ❖ Fixed-size or varying-size data structures

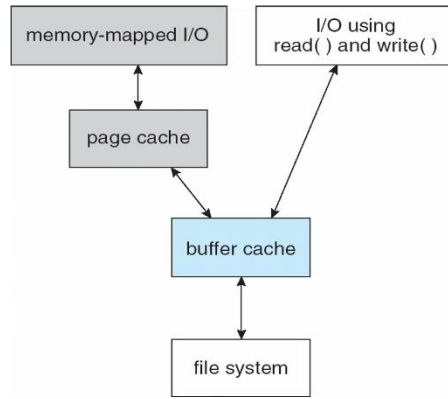
Efficiency and Performance

❖ Performance

- ❖ Keeping data and metadata close together
- ❖ **Buffer cache** – separate section of main memory for frequently used blocks
- ❖ **Synchronous** writes sometimes requested by apps or needed by OS
 - ❖ No buffering / caching - writes done on disk directly
 - ❖ **Asynchronous** writes more common, buffered, faster

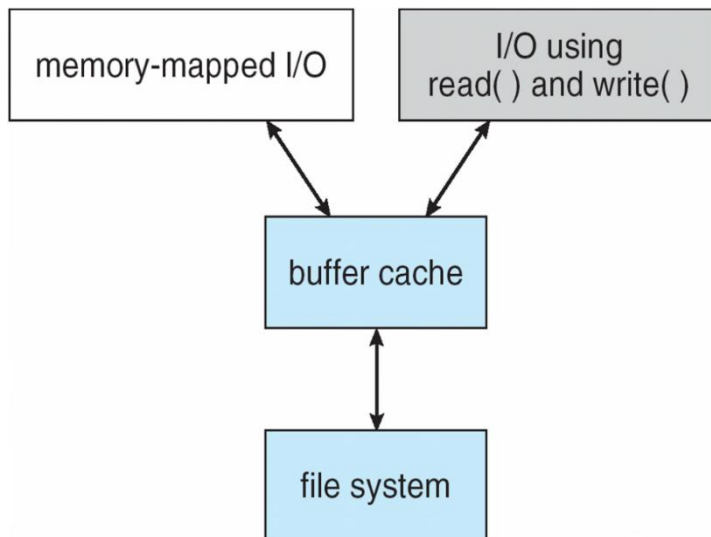
Page Cache and Buffer Cache

- ❖ A **page cache** caches pages rather than disk blocks using virtual memory techniques and addresses
- ❖ Memory-mapped I/O uses a page cache
- ❖ Routine I/O through the file system uses the buffer (disk) cache



Unified Buffer Cache

- ❖ A **unified buffer cache** uses the same page cache to cache both memory-mapped pages and ordinary file system I/O to avoid **double caching**



Recovery

- ❖ **Consistency checking** – compares data in directory structure with data blocks on disk, and tries to fix inconsistencies
 - ❖ Can be slow and sometimes fails
- ❖ Use system programs to **back up** data from disk to another storage device (magnetic tape, other magnetic disk, optical)
- ❖ Recover lost file or disk by **restoring** data from backup

Log Structured File Systems

- ❖ **Log structured** (or **journaling**) file systems record each metadata update to the file system as a **transaction**
- ❖ All transactions are written to a log
 - ❖ A transaction is considered committed once it is written to the log
 - ❖ The transactions in the log are asynchronously written to the file system structures
- ❖ If the file system crashes, all remaining transactions in the log must still be performed
- ❖ Faster recovery from crash, removes chance of inconsistency of metadata

Thank you

johnjose@iitg.ac.in

<http://www.iitg.ac.in/johnjose/>

