

CS343 - Operating Systems

Module-7C

I/O Device Classifications and Performance



Dr. John Jose

Assistant Professor

Department of Computer Science & Engineering

Indian Institute of Technology Guwahati, Assam.

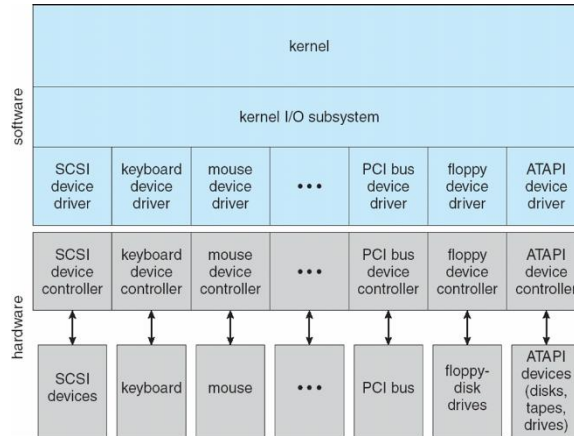
<http://www.iitg.ac.in/johnjose/>

Objectives

- ❖ Explore the structure of an operating system's I/O subsystem
- ❖ Discuss the principles of I/O hardware and its complexity
- ❖ Provide details of the performance aspects of various types of I/O devices

Application I/O Interface

- ❖ I/O system calls encapsulate device behaviors in generic classes
- ❖ Device-driver layer hides differences among I/O controllers from kernel
- ❖ New devices use already-implemented protocols. No extra work.
- ❖ Each OS has its own I/O subsystem structures and device driver frameworks



Kernel I/O Structure

Category of I/O Devices

- ❖ Devices vary in many dimensions
- ❖ **Character-stream** or **block**
- ❖ **Sequential** or **random-access**
- ❖ **Synchronous** or **asynchronous**
- ❖ **Sharable** or **dedicated**
- ❖ **Speed of operation**
- ❖ **read-write, read only, or write only**

aspect	variation	example
data-transfer mode	character block	terminal disk
access method	sequential random	modem CD-ROM
transfer schedule	synchronous asynchronous	tape keyboard
sharing	dedicated sharable	tape keyboard
device speed	latency seek time transfer rate delay between operations	
I/O direction	read only write only read-write	CD-ROM graphics controller disk

Characteristics of I/O Devices

- ❖ Broadly I/O devices can be grouped by the OS into
 - ❖ Block I/O
 - ❖ Character I/O (Stream)
 - ❖ Memory-mapped file access
 - ❖ Network sockets

Block and Character Devices

- ❖ Block devices include disk drives
 - ❖ Commands include read, write, seek
 - ❖ Raw I/O, direct I/O, or file-system access
 - ❖ File mapped to virtual memory and brought via demand paging
 - ❖ DMA
- ❖ Character devices include keyboards, mice, serial ports
 - ❖ Commands include **get()**, **put()**
 - ❖ Buffering and editing services

Network Devices

- ❖ Linux, Unix, Windows and many others include **socket** interface
 - ❖ Separates network protocol from network operation
 - ❖ Includes **select()** functionality to choose sockets for send and receive
- ❖ Pipes, FIFOs, streams, queues, mailboxes

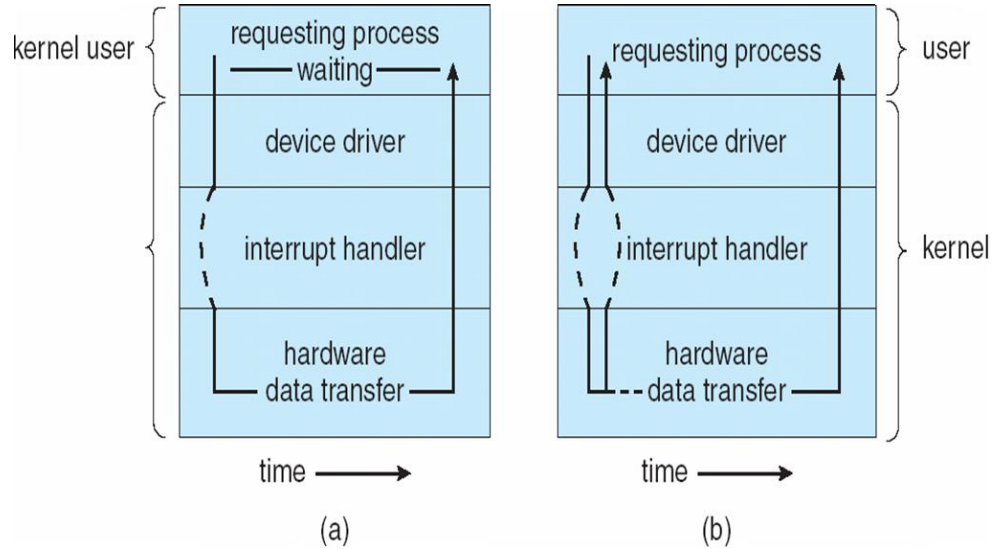
Clocks and Timers

- ❖ Provide current time, elapsed time, set a timer to trigger X and T
- ❖ Normal resolution about 1/60 second
- ❖ Some systems provide higher-resolution timers
- ❖ Programmable interval timer used for timings, periodic interrupts

Nonblocking and Asynchronous I/O

- ❖ **Blocking** - process suspended until I/O completed
 - ❖ Application from running to waiting and to ready states
- ❖ **Nonblocking** - I/O call returns as much as available I/O data
 - ❖ User interface, data copy (buffered I/O)
 - ❖ Implemented via multi-threading
 - ❖ Returns quickly with count of bytes read or written
 - ❖ **select()** to find if data ready then **read()** or **write()** to transfer
- ❖ **Asynchronous** - process runs while I/O executes
 - ❖ Difficult to use
 - ❖ I/O subsystem signals process when I/O completed

Two I/O Methods



Synchronous

Asynchronous

Kernel I/O Subsystem


❖ Scheduling

- ❖ Some I/O request ordering via per-device queue
- ❖ Some OSs try fairness
- ❖ Some implement Quality Of Service

Kernel I/O Subsystem

- ❖ **Buffering** - store data in memory while transferring between devices
 - ❖ To cope with device speed mismatch
 - ❖ To cope with device transfer size mismatch
 - ❖ To maintain copy semantics
- ❖ **Double buffering** – two copies of the data
 - ❖ Kernel and user
 - ❖ Varying sizes
 - ❖ Full / being processed and not-full / being used
 - ❖ Copy-on-write can be used for efficiency in some cases

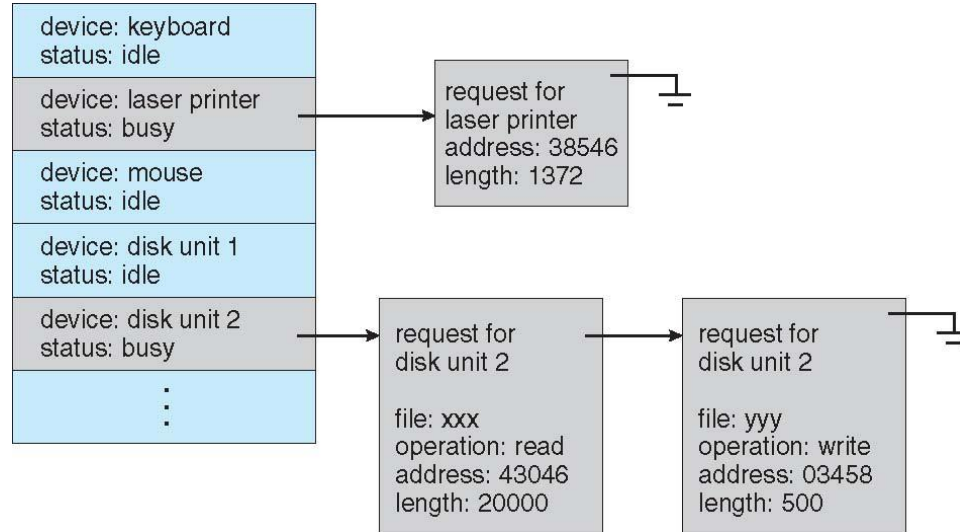
Sun Enterprise 6000 Device-Transfer Rates



Kernel I/O Subsystem

- ❖ **Caching** - faster device holding copy of data
 - ❖ Always just a copy
 - ❖ Key to performance
 - ❖ Sometimes combined with buffering
- ❖ **Spooling** - hold output for a device
 - ❖ If device can serve only one request at a time
 - ❖ i.e., Printing
- ❖ **Device reservation** - provides exclusive access to a device
 - ❖ System calls for allocation and de-allocation
 - ❖ Watch out for deadlock

Device-Status Table

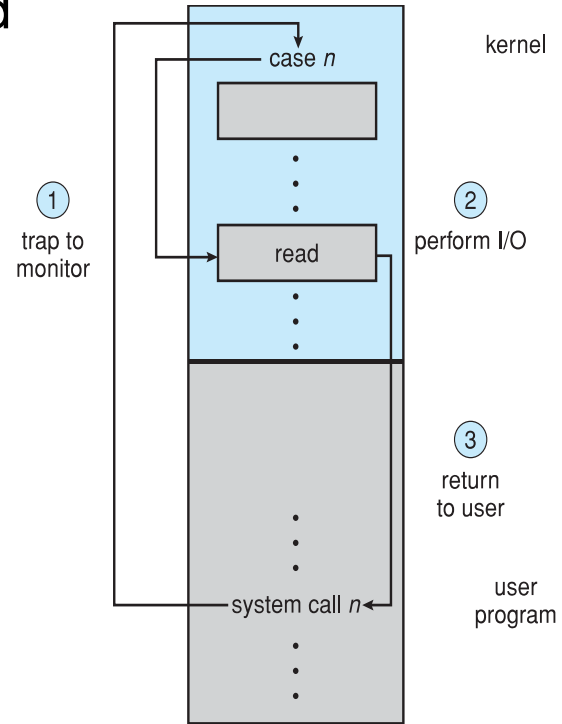


Error Handling

- ❖ OS can recover from disk read, device unavailable, transient write failures
 - ❖ Retry a read or write, for example
 - ❖ Some systems more advanced – Solaris FMA, AIX
 - ❖ Track error frequencies, stop using device with increasing frequency of retry-able errors
- ❖ Most return an error number or code when I/O request fails
- ❖ System error logs hold problem reports


System Call to Perform I/O

- ❖ All I/O instructions defined to be privileged
- ❖ I/O must be performed via system calls



Kernel Data Structures

- ❖ Kernel keeps state information for I/O components, including open file tables, network connections, character device state



Power Management

- ❖ Not strictly domain of I/O, but much is I/O related
- ❖ Computers and devices use electricity, generate heat, frequently require cooling
- ❖ OS can help manage and improve use
- ❖ Mobile computing has power management

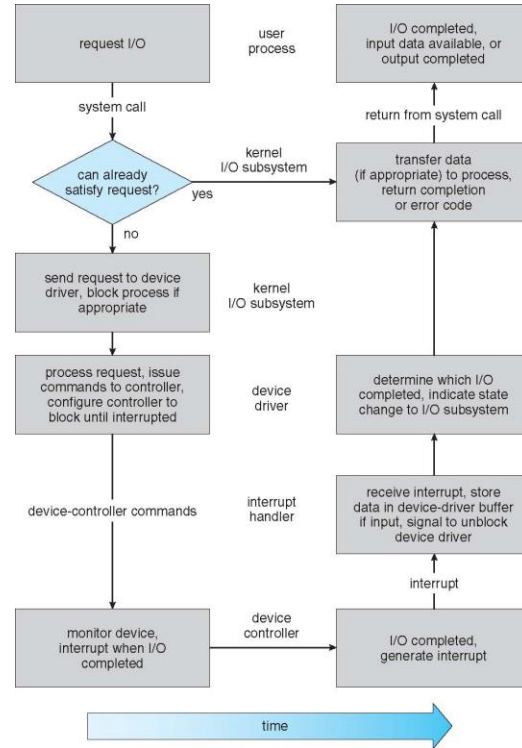
Power Management on Mobile Platforms

- ❖ Understands relationship between components
- ❖ Build device tree representing physical device topology
- ❖ System bus -> I/O subsystem -> {flash, USB storage}
- ❖ Device driver tracks state of device, whether in use
- ❖ Unused component – turn it off
- ❖ All devices in tree branch unused – turn off branch
- ❖ Wake locks – like other locks but prevent sleep of device when lock is held
- ❖ Power collapse – put a device into very deep sleep
 - ❖ Only awake enough to respond to external stimuli

I/O Requests to Hardware Operations

- ❖ Consider reading a file from disk for a process:
 - ❖ Determine device holding file
 - ❖ Translate name to device representation
 - ❖ Physically read data from disk into buffer
 - ❖ Make data available to requesting process
 - ❖ Return control to process


Life Cycle of An I/O Request



STREAMS

- ❖ **STREAM** – a full-duplex communication channel between a user-level process and a device in Unix System
- ❖ A STREAM consists of:
 - ❖ STREAM head interfaces with the user process
 - ❖ driver end interfaces with the device
 - ❖ zero or more STREAM modules between them
- ❖ Each module contains a **read queue** and a **write queue**
- ❖ Message passing is used to communicate between queues
 - ❖ **Flow control** option to indicate available or busy
- ❖ Asynchronous internally, synchronous where user process communicates with stream head

The STREAMS Structure



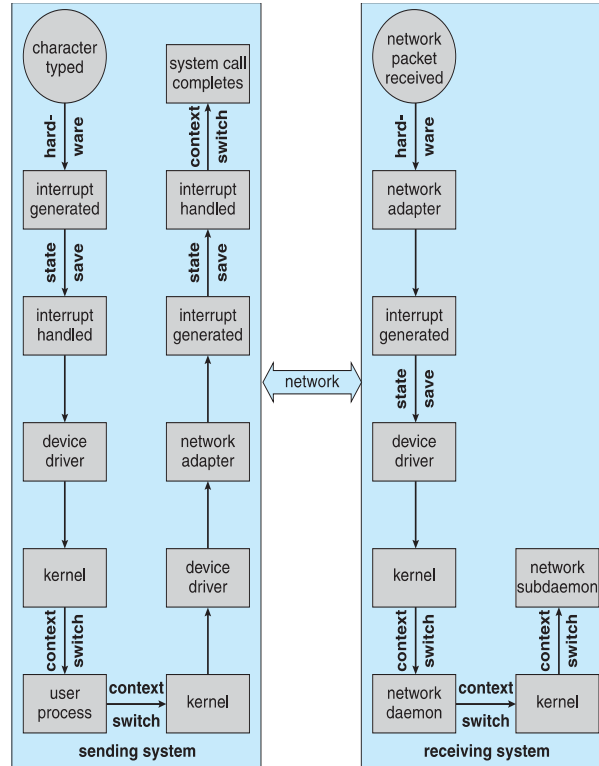
Performance

- ❖ I/O a major factor in system performance:
 - ❖ Demands CPU to execute device driver, kernel I/O code
 - ❖ Context switches due to interrupts
 - ❖ Data copying
 - ❖ Network traffic especially stressful

Improving Performance

- ❖ Reduce number of context switches
- ❖ Reduce data copying
- ❖ Reduce interrupts by using large transfers, smart controllers, polling
- ❖ Use DMA
- ❖ Use smarter hardware devices
- ❖ Balance CPU, memory, bus, and I/O performance for highest throughput
- ❖ Move user-mode processes / daemons to kernel threads

Intercomputer Communications



Thank you

johnjose@iitg.ac.in

<http://www.iitg.ac.in/johnjose/>

