

# CS343 - Operating Systems

## Module-7A

### Introduction to Input/Output Subsystem



**Dr. John Jose**

**Assistant Professor**

**Department of Computer Science & Engineering**

**Indian Institute of Technology Guwahati, Assam.**

**<http://www.iitg.ac.in/johnjose/>**

# Overview of I/O Subsystem Management

- ❖ Overview
- ❖ I/O Hardware
- ❖ Application I/O Interface
- ❖ Kernel I/O Subsystem
- ❖ Transforming I/O Requests to Hardware Operations
- ❖ I/O Performance

# Objectives

- ❖ Explore the structure of an operating system's I/O subsystem
- ❖ Discuss the principles of I/O hardware and its complexity
- ❖ Provide details of the performance aspects of I/O hardware and software

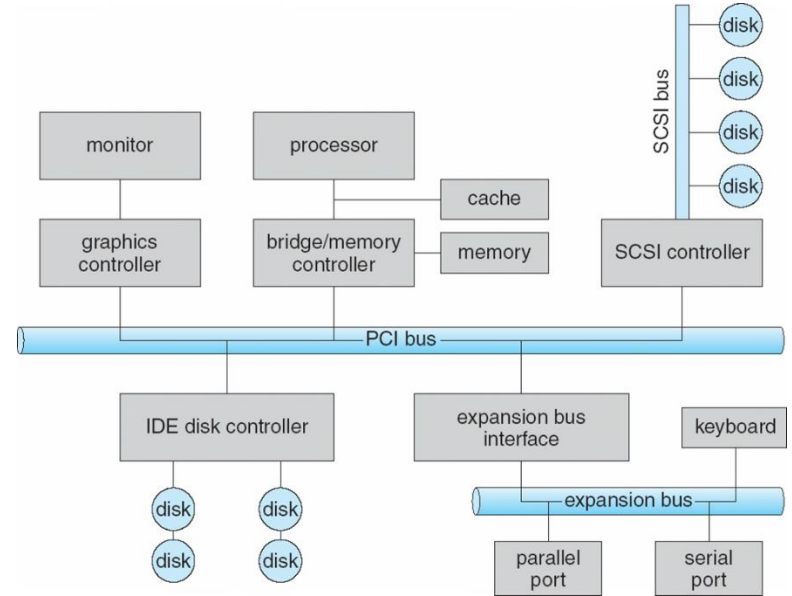
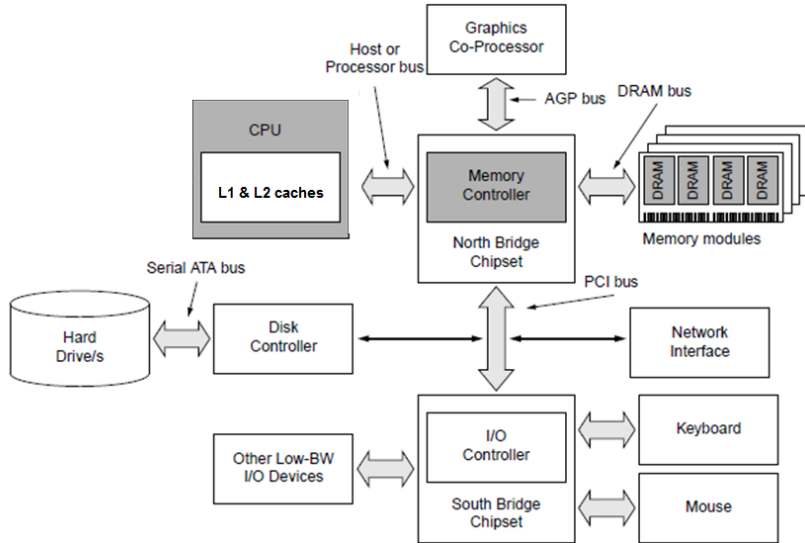
# Importance and Challenges in I/O Management

- ❖ I/O management is a major component of operating system design and operation
  - ❖ Important aspect of computer operation
  - ❖ I/O devices vary greatly
  - ❖ Various methods to control them
  - ❖ Performance management
  - ❖ New types of devices frequent
- ❖ Ports, busses, device controllers connect to various devices
- ❖ **Device drivers** encapsulate device details
  - ❖ Present uniform device-access interface to I/O subsystem

# I/O Hardware

- ❖ Incredible variety of I/O devices
  - ❖ Storage
  - ❖ Transmission
  - ❖ Human-interface

# PCI Bus Structure


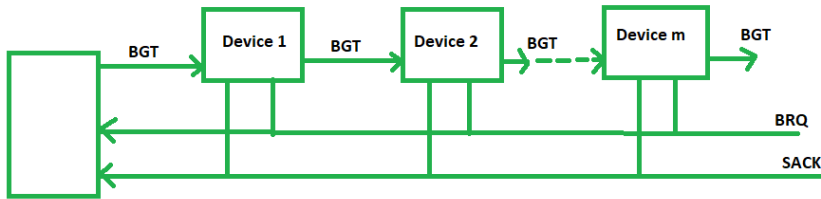


# Components of I/O Subsystem

- ❖ I/O Hardware
  - ❖ ports, buses, devices, controllers
- ❖ I/O Software
  - ❖ Interrupt Handlers, Device Driver,
  - ❖ Device-Independent Software,
  - ❖ User-Space I/O Software
- ❖ I/O Data transfer mechanisms
  - ❖ Polling, Interrupt and DMAs

# I/O Hardware

- ❖ Signals from I/O devices interface with computer
  - ❖ **Port** – connection point for device
  - ❖ **Bus** - **daisy chain** or shared direct access
    - ❖ **PCI** bus common in PCs and servers, PCI Express (**PCle**)
    - ❖ **expansion bus** connects relatively slow devices





# I/O Hardware


- ❖ **Controller** (**host adapter**) – electronics that operate port, bus, device
  - ❖ Sometimes integrated
  - ❖ Sometimes separate circuit board (host adapter)
  - ❖ Contains processor, microcode, private memory, bus controller, etc

# I/O Hardware

- ❖ I/O instructions control devices
- ❖ Devices usually have registers where device driver places commands, addresses, and data to write, or read data from registers after command execution
  - ❖ Data-in register, data-out register, status register, control register
  - ❖ Typically 1-4 bytes, or FIFO buffer
- ❖ Devices have addresses, used by I/O instructions

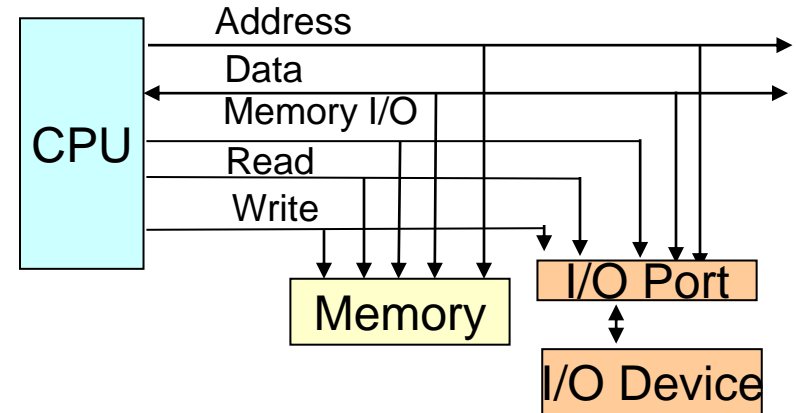

# I/O Mapping

- ❖ Memory mapped I/O
  - ❖ Devices and memory share an address space
  - ❖ I/O looks just like memory read/write
  - ❖ No special commands for I/O
  - ❖ Large selection of memory access commands available
- ❖ Isolated I/O (I/O mapped I/O)
  - ❖ Separate address spaces
  - ❖ Need I/O or memory select lines
  - ❖ Special commands for I/O; Limited set



# I/O Mapping

- ❖ CPU needs to talk to I/O
- ❖ **Memory mapped I/O**
  - ❖ Devices mapped to reserved memory locations - like RAM
  - ❖ Uses load/store instructions just like accesses to memory
- ❖ **I/O mapped I/O**
  - ❖ Special bus line
  - ❖ Special instructions



# I/O Basics

- ❖ I/O module interface I/O to CPU and Memory
- ❖ **I/O controller  $\leftrightarrow$  I/O devices ports**
  - ❖ Transfers data to/from device
  - ❖ Synchronizes operations with software
- ❖ **Status/ control registers**: device status, errors
- ❖ **Data registers**
  - ❖ Write: CPU/RAM data  $\rightarrow$  device [eg Transmit]
  - ❖ Read: CPU  $\leftarrow$  device [eg Receive]

# Functions of I/O Module

- ❖ Control & Timing
- ❖ Processor Communication
- ❖ Device Communication
- ❖ Data Buffering
- ❖ Error Detection (e.g., extra parity bit)

# Basic I/O Steps

- ❖ CPU checks I/O module device status
- ❖ I/O module returns status
- ❖ If ready, CPU requests data transfer by sending a command to the I/O module
- ❖ I/O module gets a unit of data (byte, word, etc.) from device
- ❖ I/O module transfers data to CPU
- ❖ Variations of these steps for different I/O mechanisms like polling, interrupt and DMA based I/O.

*Thank you*

**johnjose@iitg.ac.in**

**<http://www.iitg.ac.in/johnjose/>**





# CS343 - Operating Systems

## Module-7B

### Data Transfer Techniques in I/O Subsystem



**Dr. John Jose**

**Assistant Professor**

**Department of Computer Science & Engineering**

**Indian Institute of Technology Guwahati, Assam.**

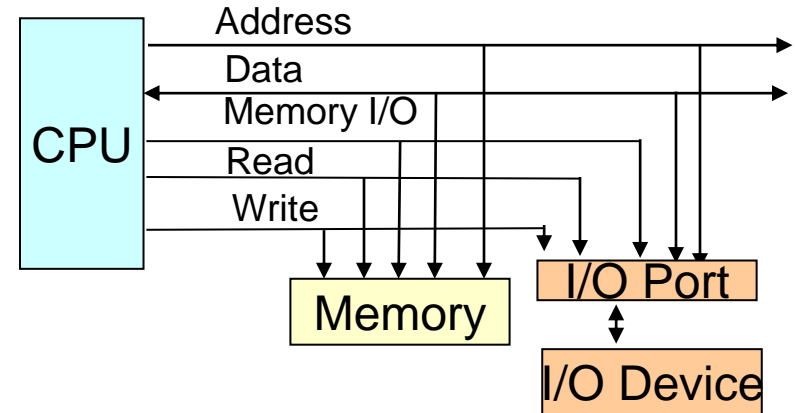
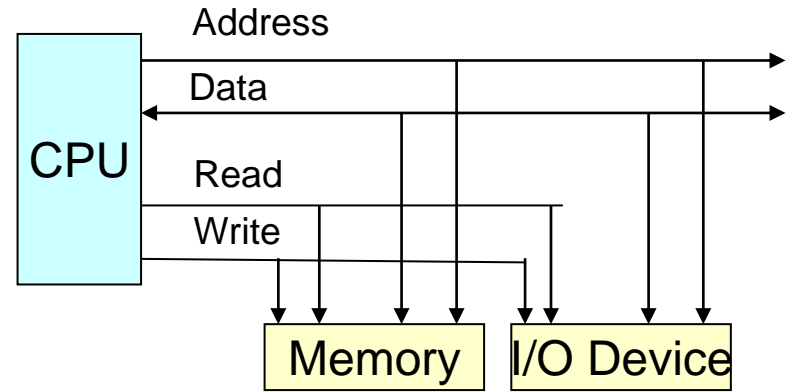
<http://www.iitg.ac.in/johnjose/>

# Objectives

- ❖ Explore the structure of an operating system's I/O subsystem
- ❖ Discuss the principles of I/O hardware and its complexity
- ❖ Provide details of the performance aspects of I/O hardware and various data transfer schemes

# I/O Mapping

- ❖ CPU needs to talk to I/O
- ❖ **Memory mapped I/O**
  - ❖ Devices mapped to reserved memory locations - like RAM
  - ❖ Uses load/store instructions just like accesses to memory
- ❖ **I/O mapped I/O**
  - ❖ Special bus line
  - ❖ Special instructions



# I/O Data Transfer techniques

- ❖ **Polled I/O**
- ❖ **Interrupt-Driven I/O**
- ❖ **Direct Memory Access (DMA)**

# Polled I/O

- ❖ CPU periodically check I/O status (polling)
  - ❖ If device ready, do operation
  - ❖ If error, take action
- ❖ CPU has direct control over I/O
  - ❖ Sensing status
  - ❖ Read/write commands
  - ❖ Transferring data
- ❖ CPU waits for I/O module to complete operation
- ❖ Wastes CPU time

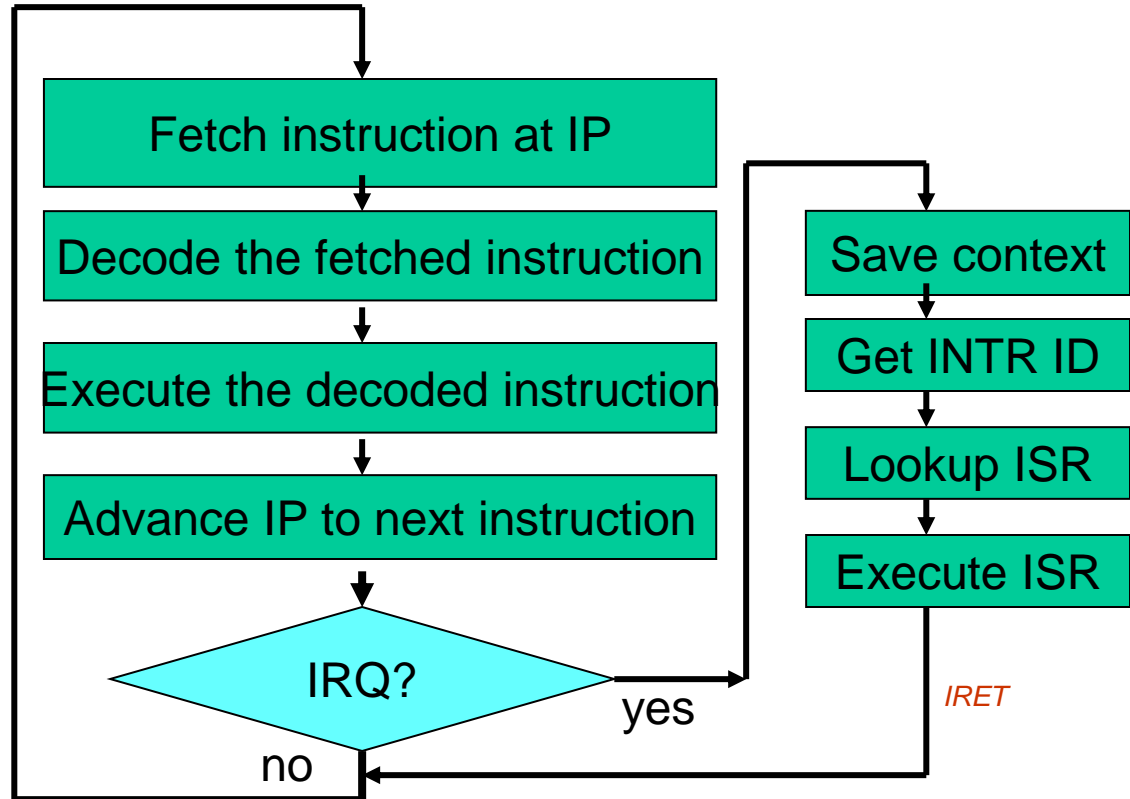
# Steps in Polled I/O

- ❖ CPU requests I/O operation
- ❖ I/O module performs operation
- ❖ I/O module sets status bits
- ❖ CPU checks status bits periodically (polling)
- ❖ CPU may wait or come back later
- ❖ I/O module does not inform CPU directly
- ❖ I/O module does not interrupt CPU

# Interrupts

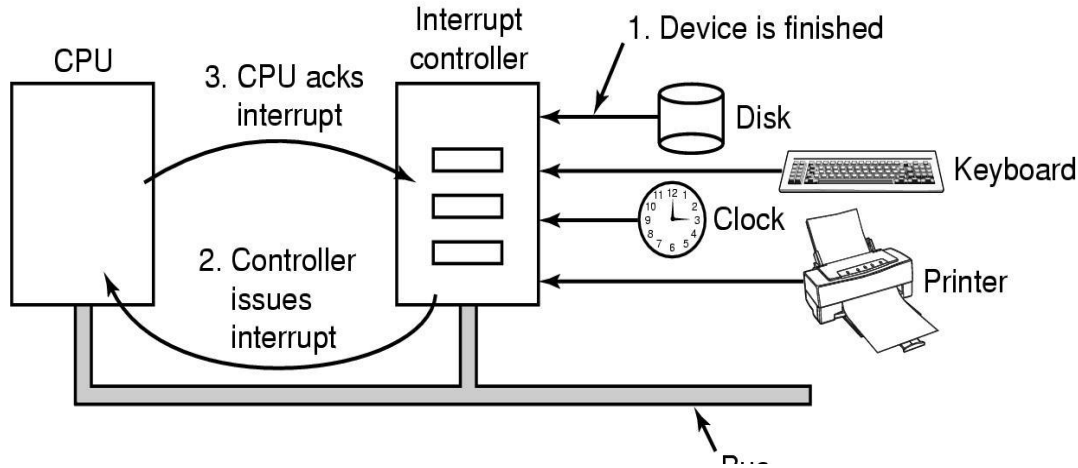
- ❖ Polling can happen in 3 instruction cycles
  - ❖ Read status, extract status bit, branch if status is shows done.
  - ❖ How to be more efficient if status is done infrequently?
- ❖ CPU **Interrupt-request line** triggered by I/O device
  - ❖ Checked by processor after each instruction
- ❖ **Interrupt handler** receives interrupts
  - ❖ **Maskable** to ignore or delay some interrupts

# Interrupt Service Routine





# Interrupt Driven I/O

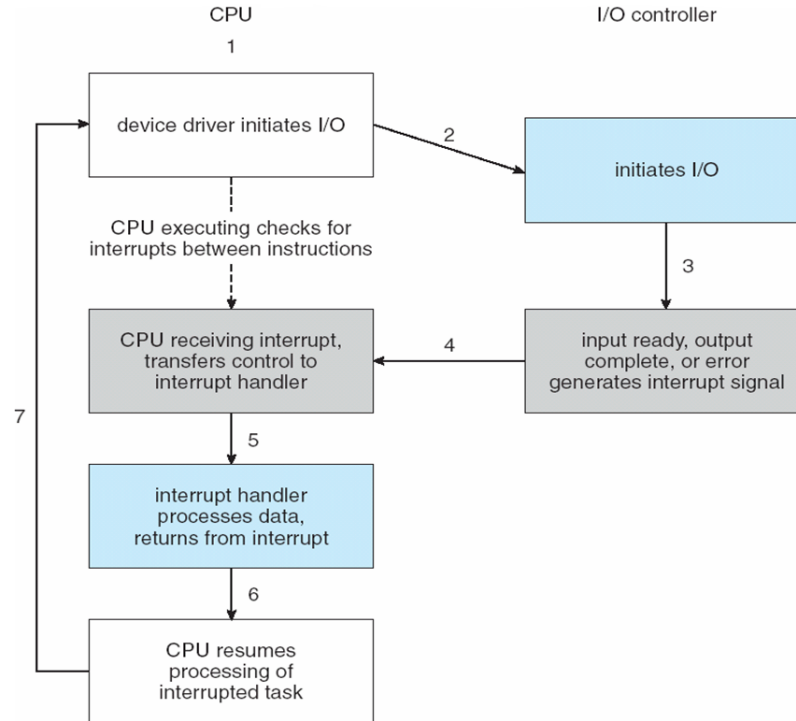


- ❖ I/O module gets data from peripheral while CPU continues other work.
- ❖ I/O module interrupts CPU when data is ready.
- ❖ CPU requests data
- ❖ I/O module transfers data

# Interrupt Driven I/O

- ❖ I/O device issues an interrupt to indicate that it needs attention of CPU
- ❖ Interrupts are special signals initiated by I/O devices to catch the attention of the processor.
- ❖ Overcomes CPU waiting
- ❖ No repeated CPU checking of device
- ❖ An I/O interrupt is **asynchronous** w.r.t. instruction execution
- ❖ Is not associated with any instruction so doesn't prevent any instruction from completing

# Interrupt-Driven I/O Cycle



# Interrupt Driven I/O

## ❖ Advantages

- ❖ Relieves the processor from having to continuously poll for an I/O event; user program progress is only suspended during the actual transfer of I/O data to/from user memory space

## ❖ Disadvantages


- ❖ Special hardware is needed to indicate the I/O device causing the interrupt and to save the necessary information prior to servicing the interrupt and to resume normal processing after servicing the interrupt

# Challenges in Interrupt Driven I/O

- ❖ How do you identify the module issuing the interrupt?
  - ❖ Need a way to identify the device generating the interrupt
- ❖ How do you deal with multiple interrupts?
  - ❖ Can have different urgencies (so need a way to prioritize them)

# Identifying Interrupting Module

- ❖ CPU asks each module in turn (Slow) Daisy Chain or Hardware poll
  - ❖ Interrupt Acknowledge sent down a chain
  - ❖ Module responsible places vector on bus
  - ❖ CPU uses vector to identify handler routine



- ❖ Vectored Interrupt
  - ❖ **Interrupt vector**  to dispatch interrupt to correct handler

# Ex: Intel Pentium Processor Event-Vector Table

vector number	description
0	divide error
1	debug exception
2	null interrupt
3	breakpoint
4	INTO-detected overflow
5	bound range exception
6	invalid opcode
7	device not available
8	double fault
9	coprocessor segment overrun (reserved)
10	invalid task state segment
11	segment not present
12	stack fault
13	general protection
14	page fault
15	(Intel reserved, do not use)
16	floating-point error
17	alignment check
18	machine check
19–31	(Intel reserved, do not use)
32–255	maskable interrupts


# Direct Memory Access

- ❖ Interrupt driven and programmed I/O require active CPU intervention
- ❖ For high-bandwidth devices (like disks) interrupt-driven I/O would consume a lot of processor cycles
- ❖ Bypasses CPU to transfer data directly between I/O device and memory
- ❖ OS writes DMA command block into memory
  - ❖ Source and destination addresses, Read or write mode
  - ❖ Count of bytes, Writes location of command block to DMA controller
- ❖ DMA is an additional module (hardware) on bus
- ❖ DMA controller takes over from CPU for I/O operations

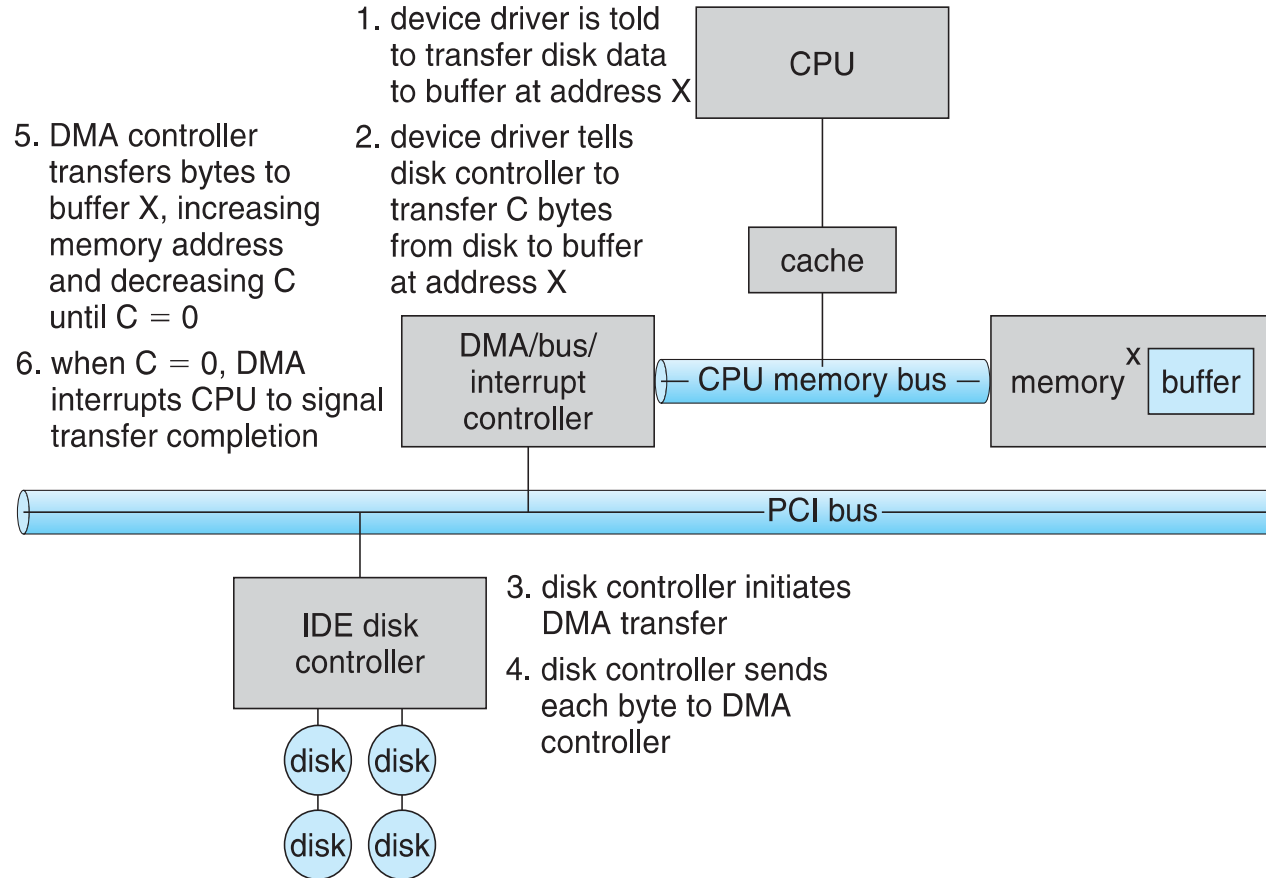


# DMA Module Diagram

- ❖ CPU tells DMA controller:-
  - ❖ Read/Write
  - ❖ Device address
  - ❖ Starting address of memory block for data
  - ❖ Amount of data to be transferred
- ❖ CPU carries on with other work
- ❖ DMA controller deals with transfer
- ❖ DMA controller sends interrupt when finished



# Disk to Memory Copy via DMA



# Modes of DMA operation

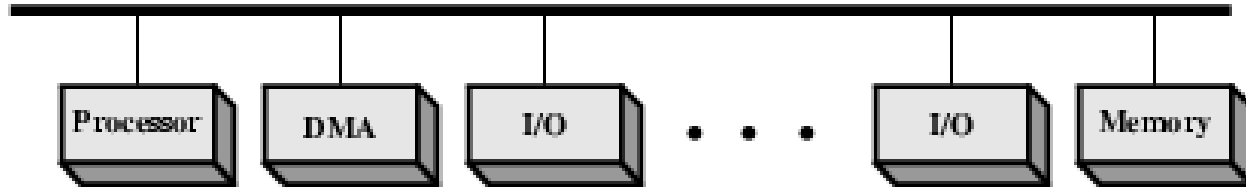
## ❖ Cycle Stealing

- ❖ DMA controller acquires control of bus
- ❖ Transfers a single word and releases the bus
- ❖ The CPU is slowed down due to bus contention
- ❖ Responsive but not very efficient

## ❖ Burst Mode


- ❖ DMA Controller acquires control of bus
- ❖ Transfers all the data and then only releases the bus
- ❖ The CPU is suspended or works with cache
- ❖ Efficient but interrupts may not be serviced in a timely way

# DMA Configurations



- ❖ Single Bus, Detached DMA controller
- ❖ Each transfer uses bus twice
  - ❖ I/O to DMA then DMA to memory
- ❖ CPU is suspended twice


# DMA Configurations



(b) **Single-bus, Integrated DMA-I/O**

- ❖ Single Bus, integrated DMA controller
- ❖ Controller may support >1 device
- ❖ Each transfer uses bus once
  - ❖ DMA to memory
- ❖ CPU is suspended once

# DMA Configurations



- ❖ Separate I/O Bus
- ❖ Bus supports all DMA enabled devices
- ❖ Each transfer uses bus once
  - ❖ DMA to memory
- ❖ CPU is suspended once

*Thank you*

**johnjose@iitg.ac.in**

**<http://www.iitg.ac.in/johnjose/>**



# CS343 - Operating Systems

## Module-7C

### I/O Device Classifications and Performance



**Dr. John Jose**

**Assistant Professor**

**Department of Computer Science & Engineering**

**Indian Institute of Technology Guwahati, Assam.**

<http://www.iitg.ac.in/johnjose/>




# Objectives

- ❖ Explore the structure of an operating system's I/O subsystem
- ❖ Discuss the principles of I/O hardware and its complexity
- ❖ Provide details of the performance aspects of various types of I/O devices

# Application I/O Interface

- ❖ I/O system calls encapsulate device behaviors in generic classes
- ❖ Device-driver layer hides differences among I/O controllers from kernel
- ❖ New devices use already-implemented protocols. No extra work.
- ❖ Each OS has its own I/O subsystem structures and device driver frameworks



Kernel I/O Structure

# Category of I/O Devices

- ❖ Devices vary in many dimensions
- ❖ **Character-stream** or **block**
- ❖ **Sequential** or **random-access**
- ❖ **Synchronous** or **asynchronous**
- ❖ **Sharable** or **dedicated**
- ❖ **Speed of operation**
- ❖ **read-write, read only, or write only**

aspect	variation	example
data-transfer mode	character block	terminal disk
access method	sequential random	modem CD-ROM
transfer schedule	synchronous asynchronous	tape keyboard
sharing	dedicated sharable	tape keyboard
device speed	latency seek time transfer rate delay between operations	
I/O direction	read only write only read-write	CD-ROM graphics controller disk

# Characteristics of I/O Devices

- ❖ Broadly I/O devices can be grouped by the OS into
  - ❖ Block I/O
  - ❖ Character I/O (Stream)
  - ❖ Memory-mapped file access
  - ❖ Network sockets

# Block and Character Devices

- ❖ Block devices include disk drives
  - ❖ Commands include read, write, seek
  - ❖ Raw I/O, direct I/O, or file-system access
  - ❖ File mapped to virtual memory and brought via demand paging
  - ❖ DMA
- ❖ Character devices include keyboards, mice, serial ports
  - ❖ Commands include **get()**, **put()**
  - ❖ Buffering and editing services

# Network Devices

- ❖ Linux, Unix, Windows and many others include **socket** interface
  - ❖ Separates network protocol from network operation
  - ❖ Includes **select()** functionality to choose sockets for send and receive
- ❖ Pipes, FIFOs, streams, queues, mailboxes

# Clocks and Timers


- ❖ Provide current time, elapsed time, set a timer to trigger  $X$  and  $T$
- ❖ Normal resolution about 1/60 second
- ❖ Some systems provide higher-resolution timers
- ❖ Programmable interval timer used for timings, periodic interrupts

# Nonblocking and Asynchronous I/O

- ❖ **Blocking** - process suspended until I/O completed
  - ❖ Application from running to waiting and to ready states
- ❖ **Nonblocking** - I/O call returns as much as available I/O data
  - ❖ User interface, data copy (buffered I/O)
  - ❖ Implemented via multi-threading
  - ❖ Returns quickly with count of bytes read or written
  - ❖ **select()** to find if data ready then **read()** or **write()** to transfer
- ❖ **Asynchronous** - process runs while I/O executes
  - ❖ Difficult to use
  - ❖ I/O subsystem signals process when I/O completed



# Two I/O Methods



Synchronous

Asynchronous

# Kernel I/O Subsystem


## ❖ Scheduling

- ❖ Some I/O request ordering via per-device queue
- ❖ Some OSs try fairness
- ❖ Some implement Quality Of Service

# Kernel I/O Subsystem

- ❖ **Buffering** - store data in memory while transferring between devices
  - ❖ To cope with device speed mismatch
  - ❖ To cope with device transfer size mismatch
  - ❖ To maintain copy semantics
- ❖ **Double buffering** – two copies of the data
  - ❖ Kernel and user
  - ❖ Varying sizes
  - ❖ Full / being processed and not-full / being used
  - ❖ Copy-on-write can be used for efficiency in some cases


# Sun Enterprise 6000 Device-Transfer Rates



# Kernel I/O Subsystem

- ❖ **Caching** - faster device holding copy of data
  - ❖ Always just a copy
  - ❖ Key to performance
  - ❖ Sometimes combined with buffering
- ❖ **Spooling** - hold output for a device
  - ❖ If device can serve only one request at a time
  - ❖ i.e., Printing
- ❖ **Device reservation** - provides exclusive access to a device
  - ❖ System calls for allocation and de-allocation
  - ❖ Watch out for deadlock

# Device-Status Table




# Error Handling

- ❖ OS can recover from disk read, device unavailable, transient write failures
  - ❖ Retry a read or write, for example
  - ❖ Some systems more advanced – Solaris FMA, AIX
    - ❖ Track error frequencies, stop using device with increasing frequency of retry-able errors
- ❖ Most return an error number or code when I/O request fails
- ❖ System error logs hold problem reports

# System Call to Perform I/O

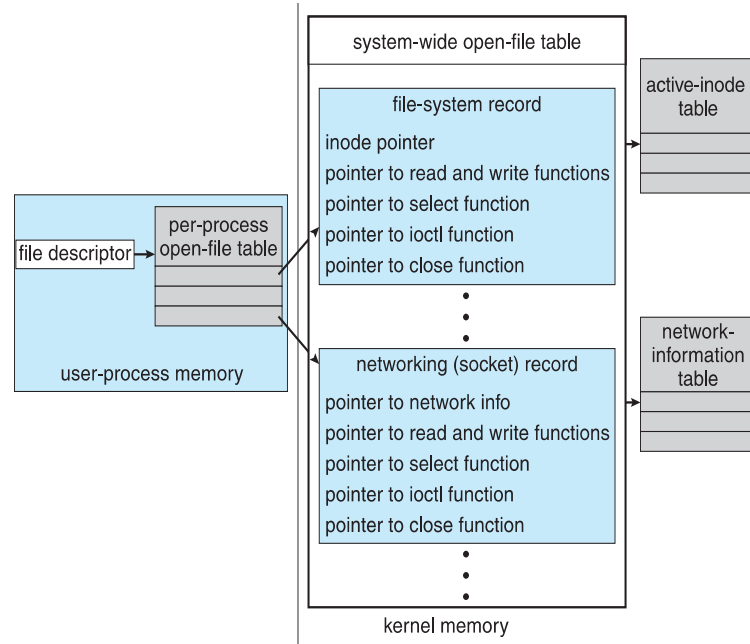
- ❖ All I/O instructions defined to be privileged
- ❖ I/O must be performed via system calls





# Kernel Data Structures

- ❖ Kernel keeps state information for I/O components, including open file tables, network connections, character device state



# Power Management

- ❖ Not strictly domain of I/O, but much is I/O related
- ❖ Computers and devices use electricity, generate heat, frequently require cooling
- ❖ OS can help manage and improve use
- ❖ Mobile computing has power management


# Power Management on Mobile Platforms

- ❖ Understands relationship between components
- ❖ Build device tree representing physical device topology
- ❖ System bus -> I/O subsystem -> {flash, USB storage}
- ❖ Device driver tracks state of device, whether in use
- ❖ Unused component – turn it off
- ❖ All devices in tree branch unused – turn off branch
- ❖ Wake locks – like other locks but prevent sleep of device when lock is held
- ❖ Power collapse – put a device into very deep sleep
  - ❖ Only awake enough to respond to external stimuli

# I/O Requests to Hardware Operations

- ❖ Consider reading a file from disk for a process:
  - ❖ Determine device holding file
  - ❖ Translate name to device representation
  - ❖ Physically read data from disk into buffer
  - ❖ Make data available to requesting process
  - ❖ Return control to process


# Life Cycle of An I/O Request



# STREAMS

- ❖ **STREAM** – a full-duplex communication channel between a user-level process and a device in Unix System
- ❖ A STREAM consists of:
  - ❖ STREAM head interfaces with the user process
  - ❖ driver end interfaces with the device
  - ❖ zero or more STREAM modules between them
- ❖ Each module contains a **read queue** and a **write queue**
- ❖ Message passing is used to communicate between queues
  - ❖ **Flow control** option to indicate available or busy
- ❖ Asynchronous internally, synchronous where user process communicates with stream head

# The STREAMS Structure



# Performance


- ❖ I/O a major factor in system performance:
  - ❖ Demands CPU to execute device driver, kernel I/O code
  - ❖ Context switches due to interrupts
  - ❖ Data copying
  - ❖ Network traffic especially stressful



# Improving Performance

- ❖ Reduce number of context switches
- ❖ Reduce data copying
- ❖ Reduce interrupts by using large transfers, smart controllers, polling
- ❖ Use DMA
- ❖ Use smarter hardware devices
- ❖ Balance CPU, memory, bus, and I/O performance for highest throughput
- ❖ Move user-mode processes / daemons to kernel threads

# Intercomputer Communications



*Thank you*

**johnjose@iitg.ac.in**

**<http://www.iitg.ac.in/johnjose/>**

